

Importing the libraries needed

In [1]:

```
import numpy as np
import pandas as pd
import time

import matplotlib.pyplot as plt
import seaborn as sns

import re
import string

from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score

import gensim
from gensim.models import KeyedVectors

from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences

import tensorflow as tf
from keras.models import Sequential
from tensorflow.keras.layers import SpatialDropout1D, Conv1D, Bidirectional, LSTM, Dense, Input, Dropout, GlobalMaxPooling1D
from keras.layers.embeddings import Embedding
from tensorflow.keras.callbacks import ModelCheckpoint, ReduceLROnPlateau, EarlyStopping
from tensorflow.keras.optimizers import Adam

import itertools
from numpy import loadtxt
from keras.models import load_model

import warnings
warnings.filterwarnings("ignore")
```

Connecting to google drive

In [2]:

```
from google.colab import drive
drive.mount("/content/gdrive")
```

Mounted at /content/gdrive

Uploading the dataset

In [3]:

```
path_data = "/content/gdrive/MyDrive/thesis/modified.csv"
Arsas = pd.read_csv(path_data ,sep='\t')
```

In [4]:

```
data = Arsas
```

printing the first 3 rows of the data

In [5]:

```
data.head(3)
```

Out[5]:

	#Tweet_ID	Tweet_text	Sentiment_label
0	929241870508724224	مصر الجولة الأخيرة x المباراة القادمة #غانا	Positive
1	928942264583376897	هل هذه هي سياسة خارجيه لدوله تحترم نفسها والآخ	Negative
2	928615163250520065	وزير خارجية فرنسا عن منتدى شباب العالم: شعرت ب	Positive

printing the shape of the dataset nbr of row and columns

In [6]:

```
print("Data contient {} lignes et {} colonnes.".format(data.shape[0], data.shape[1]))
```

Data contient 21064 lignes et 3 colonnes.

printing the fiels with missed values

In [7]:

```
data.isnull().sum()
```

Out[7]:

```
#Tweet_ID      0
Tweet_text      0
Sentiment_label 0
dtype: int64
```

printing the number of the duplicated rows

In [8]:

```
print("On a {} doublons dans Data.".format(data.duplicated().sum()))
```

On a 68 doublons dans Data.

In [9]:

```
data.drop_duplicates(inplace = True)
```

In [10]:

```
print("On a {} doublons dans Data.".format(data.duplicated().sum()))
```

On a 0 doublons dans Data.

checking the types of the fiels in the data

In [11]:

```
data.dtypes
```

Out[11]:

```
#Tweet_ID          int64
Tweet_text          object
Sentiment_label     object
dtype: object
```

function for printing the pie

In [12]:

```
def pie(data,col):
    labels = data[col].value_counts().keys().tolist()
    n = len(labels)
    if n==2:
        colors = ['#66b3ff', '#fb3999']
    elif n==3:
        colors = ['#66b3ff', '#fb3999', '#ffcc99']
    elif n==4:
        colors = ['#66b3ff', '#fb3999', '#ffcc99', "#66f3ff"]
    elif n==5:
        colors = ['#66b3ff', '#fb3999', '#ffcc99', "#66f3ff", '#adcc99']
    elif n==6:
        colors = ['#66b3ff', '#fb3999', '#ffcc99', "#66f3ff", '#adcc99', "#db7f23"]

    fig1, f1 = plt.subplots()
    f1.pie(data[col].value_counts(), labels=labels, colors = colors, autopct='%
1.1f%%',shadow=False, startangle=60)
    f1.axis('equal')
    plt.tight_layout()
    plt.show()

def histo(data,col):
    plt.figure(figsize = (10, 8))
    sns.histplot(data=data, x=col, hue = data[col], fill=True)
```

Counting the % of each classe

In [13]:

```
data.Sentiment_label.value_counts(normalize = True)
```


Out[13]:

```
Negative    0.371404  
Neutral     0.346018  
Positive    0.220566  
Mixed       0.062012  
Name: Sentiment_label, dtype: float64
```

Printing the distribution of the classes

In [14]:

```
pie(data, "Sentiment_label")
```



In [15]:

```
positive = data[data["Sentiment_label"] == "Positive"]  
positive["sentiment"] = 1  
  
mixed = data[data["Sentiment_label"] == "Mixed"]  
mixed["sentiment"] = 2  
  
neutral = data[data["Sentiment_label"] == "Neutral"]  
neutral["sentiment"] = 3  
  
negative = data[data["Sentiment_label"] == "Negative"]  
negative["sentiment"] = 0  
  
data = pd.concat([positive, mixed, neutral, negative], ignore_index = True)
```

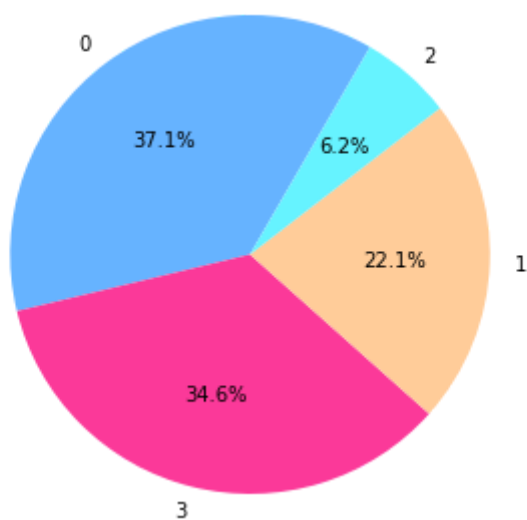
In [16]:

```
print("data contient {} lignes.".format(data.shape[0]))  
print("Positive contient {} lignes.".format(positive.shape[0]))  
print("Negative contient {} lignes.".format(negative.shape[0]))  
print("Mixed contient {} lignes.".format(mixed.shape[0]))  
print("Neutral contient {} lignes.".format(neutral.shape[0]))
```

data contient 20996 lignes.
Positive contient 4631 lignes.
Negative contient 7798 lignes.
Mixed contient 1302 lignes.
Neutral contient 7265 lignes.

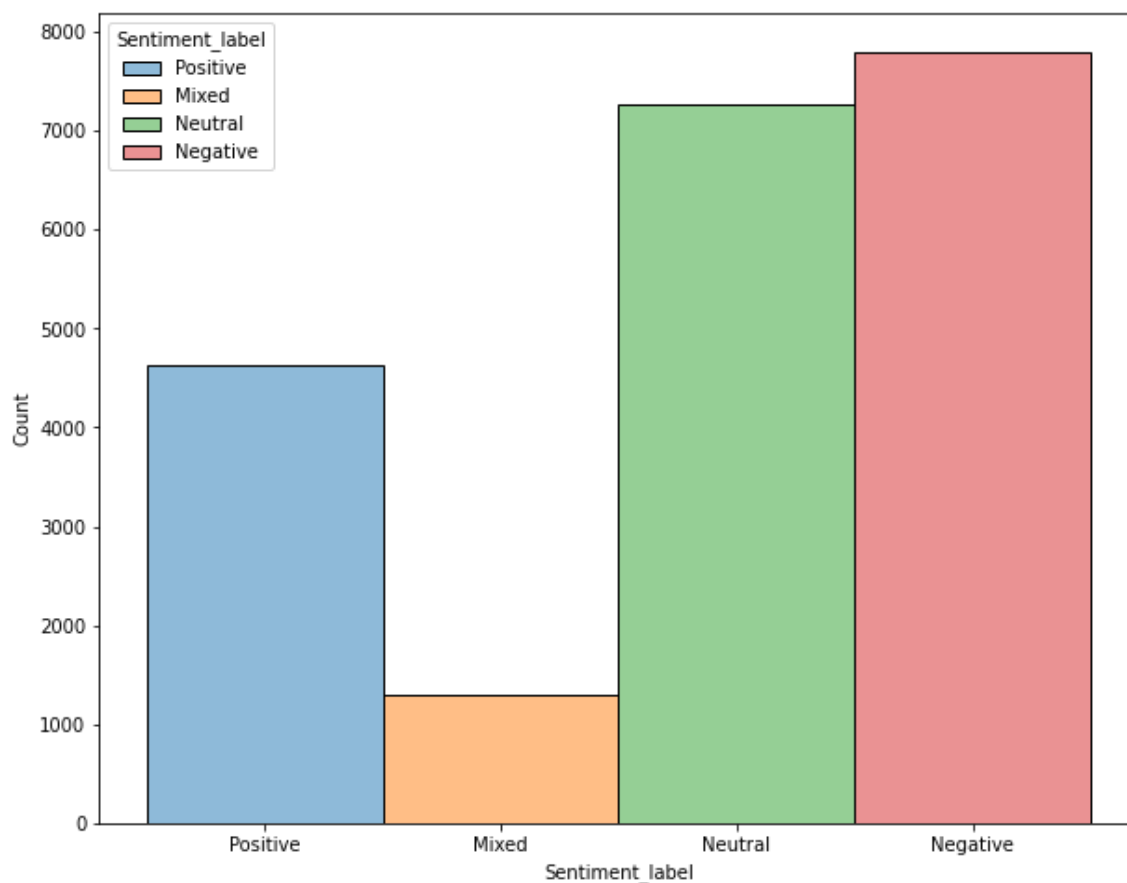
In [17]:

```
pie(data, "sentiment")
```



In [18]:

```
histo(data, "Sentiment_label")
```



function to count the length of reviews

In [19]:

```
def compte_mots(phrase):  
    return len(phrase.split())
```

```
data["len_review"] = data["Tweet_text"].apply(compte_mots)
```

printing the max length of the positive and negative reviews

In [20]:

```
print("Le maximum de mots utilisé dans les reviews est :", max(data['len_review']))  
print("Le moyen de mots utilisé dans les reviews est :", np.mean(data['len_review']))
```

Le maximum de mots utilisé dans les reviews est : 64


Le moyen de mots utilisé dans les reviews est : 19.701657458563535

In [21]:

```
plt.figure(figsize=(10,9))  
  
pl=sns.kdeplot(data['len_review'], hue = data['Sentiment_label'], shade=True, color="r")  
  
plt.xlim(0, 100)
```

Out[21]:

(0.0, 100.0)



Deleting unused fields

In [22]:

```
data.drop(['#Tweet_ID'], axis = 1, inplace = True)
data.head(3)
```

Out[22]:

	Tweet_text	Sentiment_label	sentiment	len_review
0	... مصر الجولة الأخيرة x المباراة القادمة #غانا	Positive	1	45
1	...وزير خارجية فرنسا عن منتدى شباب العالم: شعرت ب	Positive	1	16
2	بسم الله نبداً 🌟 نغرد علي وسم 👍 👍 👍 👍 👍 #شباب...	Positive	1	27

In [23]:

```
df = data
df.dtypes
```

Out[23]:

```
Tweet_text      object
Sentiment_label  object
sentiment        int64
len_review       int64
dtype: object
```

the function of the preprocessing

In [24]:

```
def preprocessing(text):

    # ref: https://github.com/bakrianoo/aravec
    search = ["ى", "ي", "-", "'", "ل", "و", ":", "/", "-_", "ة", "آ", "إ", "أ", "\\", '\n', '\t', '&quot;', '?', '¿', '!']
    replace = ["ل", "و", "", "", "", "", "", "ه", "|", "|", |", """, "", "", "?", ' ', ' ', ' ', ' ', ' ', ' ', "ي", ' ', ' ! ' ]

    tashkeel = re.compile(r'[\u0617-\u061A\u064B-\u0652]')
    text = re.sub(tashkeel, "", text)

    longation = re.compile(r'(\.)\1+')
    subst = r"\1\1"
    text = re.sub(longation, subst, text)

    text = re.sub(r"^[^w\s]", '', text)
    text = re.sub(r"[a-zA-Z]", '', text)
    text = re.sub(r"\d+", '', text)
    text = re.sub(r"\n+", '', text)
    text = re.sub(r"\t+", '', text)
    text = re.sub(r"\r+", '', text)
    text = re.sub(r"s+", '', text)
    text = text.replace('و', 'وو')
    text = text.replace('ي', 'يي')
    text = text.replace('|', '||')

    for i in range(0, len(search)):
        text = text.replace(search[i], replace[i])

    text = text.strip()

    return text
```

preprocessing the reviews and printing the time spent

In [25]:

```
%%time
data["Clean reviews"] = data.Tweet_text.apply(lambda x: preprocessing(x))
```

```
CPU times: user 1.03 s, sys: 4.96 ms, total: 1.03 s
Wall time: 1.04 s
```

printing a review before and after preprocessing

In [26]:

```
print('- Avant le prétraitement \n\n',data["Tweet_text"][4])
print("\n-----\n")
print('- Après le prétraitement \n\n',data["Clean_reviews"][4])
```

- Avant le prétraitement

http://t.co/t1Q0l0U1P لدعم محمد صلاح للحصول على جائزة الأفضل بأفريقيا «BBC» شارك بتصويت

- Après le prétraitement

شارك بتصويت لدعم محمد صلاح للحصول على جائزة الأفضل بأفريقيا

Saving the cleaned data in a csv file

In [27]:

```
data.to_csv("cleaned_Arsas.csv")
```

assigning the reviews and classes to a new variables

In [28]:

```
X = data.Clean_reviews
y=pd.get_dummies(data.sentiment)
# y = data.sentiment
```

splitting the data to train and test set

In [29]:

```
X_train, X_test, y_train, y_test = train_test_split(X,
                                                    y,
                                                    test_size = 0.20,
                                                    random_state = 42)
```

printing the number of the train set and the test set

In [30]:

```
print('Train set', X_train.shape)
print('Test set', X_test.shape)
```

Train set (16796,)
Test set (4200,)

In [31]:

```
from google.colab import drive
drive.mount('/content/gdrive')
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True).

Uploading the fsttext pretrained word embedding with 150 dimension

In [32]:

```
%%time
target_word_vec = KeyedVectors.load_word2vec_format("/content/gdrive/MyDrive/the
sis/cc.ar.150.vec", binary = False)
```

CPU times: user 2min 43s, sys: 4.03 s, total: 2min 47s
Wall time: 3min 5s

tokenization of the reviews

In [33]:

```
%%time
tokenizer = Tokenizer()
tokenizer.fit_on_texts(X_train)
```

CPU times: user 499 ms, sys: 11.1 ms, total: 510 ms
Wall time: 510 ms

In [34]:

```
word_index = tokenizer.word_index
vocab_size = len(tokenizer.word_index) + 1
```

making all reviews of the same length 70

In [35]:

```
%%time
MAX_SEQUENCE_LENGTH = 70

X_train = pad_sequences(tokenizer.texts_to_sequences(X_train),
                        maxlen = MAX_SEQUENCE_LENGTH)
X_test = pad_sequences(tokenizer.texts_to_sequences(X_test),
                      maxlen = MAX_SEQUENCE_LENGTH)

print("Training X Shape:", X_train.shape)
print("Testing X Shape:", X_test.shape)
```

Training X Shape: (16796, 70)
Testing X Shape: (4200, 70)
CPU times: user 529 ms, sys: 6.09 ms, total: 536 ms
Wall time: 536 ms

Construction of the embedding matrix

In [36]:

```
%%time
embedding_matrix = np.zeros((vocab_size, 150))

for word, i in word_index.items():
    if word in target_word_vec :
        embedding_vector = target_word_vec[word]
        if embedding_vector is not None:
            embedding_matrix[i] = embedding_vector
```

CPU times: user 172 ms, sys: 20 ms, total: 192 ms
Wall time: 192 ms

In [37]:

```
embedding_matrix.shape[0] == vocab_size
```

Out[37]:

True

Creating the model

In [38]:

```
model = Sequential()
embedding_layer = Embedding(vocab_size,
                             150,
                             weights = [embedding_matrix],
                             input_length = MAX_SEQUENCE_LENGTH,
                             trainable=False)
model.add(embedding_layer)
model.add(Conv1D(filters=64, kernel_size=2, activation='relu'))
model.add(Bidirectional(LSTM(64, dropout=0.2, return_sequences=True)))
model.add(GlobalMaxPooling1D())
model.add(Dropout(0.2))
model.add(Dense(4, activation='softmax'))

model.compile(optimizer = Adam(learning_rate=0.001),
              loss = 'categorical_crossentropy',
              metrics = ['accuracy'])

es = EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=5)
print(model.summary())
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 70, 150)	7419000
conv1d (Conv1D)	(None, 69, 64)	19264
bidirectional (Bidirectional)	(None, 69, 128)	66048
global_max_pooling1d (GlobalMaxPooling1D)	(None, 128)	0
dropout (Dropout)	(None, 128)	0
dense (Dense)	(None, 4)	516
Total params: 7,504,828		
Trainable params: 85,828		
Non-trainable params: 7,419,000		
None		

fitting the model to the dataset

In [39]:


```
history = model.fit(X_train, y_train, validation_split=0.15, batch_size = 128, epochs=20, verbose=1, callbacks=[es])
```

Epoch 1/20
112/112 [=====] - 18s 23ms/step - loss: 1.0
187 - accuracy: 0.5796 - val_loss: 0.8690 - val_accuracy: 0.6532
Epoch 2/20
112/112 [=====] - 2s 18ms/step - loss: 0.81
44 - accuracy: 0.6884 - val_loss: 0.8012 - val_accuracy: 0.6988
Epoch 3/20
112/112 [=====] - 2s 18ms/step - loss: 0.74
96 - accuracy: 0.7198 - val_loss: 0.7262 - val_accuracy: 0.7206
Epoch 4/20
112/112 [=====] - 2s 20ms/step - loss: 0.72
79 - accuracy: 0.7282 - val_loss: 0.7135 - val_accuracy: 0.7270
Epoch 5/20
112/112 [=====] - 1s 13ms/step - loss: 0.71
30 - accuracy: 0.7335 - val_loss: 0.7165 - val_accuracy: 0.7290
Epoch 6/20
112/112 [=====] - 1s 13ms/step - loss: 0.69
63 - accuracy: 0.7422 - val_loss: 0.7090 - val_accuracy: 0.7270
Epoch 7/20
112/112 [=====] - 1s 13ms/step - loss: 0.68
44 - accuracy: 0.7487 - val_loss: 0.7025 - val_accuracy: 0.7294
Epoch 8/20
112/112 [=====] - 1s 13ms/step - loss: 0.67
37 - accuracy: 0.7481 - val_loss: 0.6882 - val_accuracy: 0.7337
Epoch 9/20
112/112 [=====] - 1s 13ms/step - loss: 0.66
35 - accuracy: 0.7560 - val_loss: 0.6903 - val_accuracy: 0.7409
Epoch 10/20
112/112 [=====] - 1s 13ms/step - loss: 0.65
26 - accuracy: 0.7572 - val_loss: 0.6903 - val_accuracy: 0.7381
Epoch 11/20
112/112 [=====] - 1s 13ms/step - loss: 0.64
46 - accuracy: 0.7613 - val_loss: 0.6938 - val_accuracy: 0.7357
Epoch 12/20
112/112 [=====] - 1s 13ms/step - loss: 0.63
89 - accuracy: 0.7630 - val_loss: 0.6832 - val_accuracy: 0.7361
Epoch 13/20
112/112 [=====] - 1s 13ms/step - loss: 0.62
86 - accuracy: 0.7686 - val_loss: 0.6989 - val_accuracy: 0.7206
Epoch 14/20
112/112 [=====] - 1s 13ms/step - loss: 0.61
75 - accuracy: 0.7754 - val_loss: 0.6916 - val_accuracy: 0.7317
Epoch 15/20
112/112 [=====] - 1s 13ms/step - loss: 0.60
86 - accuracy: 0.7744 - val_loss: 0.6817 - val_accuracy: 0.7385
Epoch 16/20
112/112 [=====] - 1s 13ms/step - loss: 0.60
70 - accuracy: 0.7770 - val_loss: 0.6873 - val_accuracy: 0.7313
Epoch 17/20
112/112 [=====] - 2s 16ms/step - loss: 0.59
59 - accuracy: 0.7784 - val_loss: 0.6890 - val_accuracy: 0.7361
Epoch 18/20
112/112 [=====] - 1s 13ms/step - loss: 0.59
36 - accuracy: 0.7810 - val_loss: 0.6872 - val_accuracy: 0.7333
Epoch 19/20
112/112 [=====] - 1s 13ms/step - loss: 0.58
73 - accuracy: 0.7820 - val_loss: 0.6795 - val_accuracy: 0.7373
Epoch 20/20
112/112 [=====] - 1s 13ms/step - loss: 0.56
57 - accuracy: 0.7894 - val_loss: 0.6885 - val_accuracy: 0.7333

Evaluating the model


In [40]:

```
loss_train = history.history['loss']
loss_val = history.history['val_loss']
epochs = range(1,21)
plt.plot(epochs, loss_train, 'g', label='Training loss')
plt.plot(epochs, loss_val, 'b', label='validation loss')
plt.title('Training and Validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



In [41]:

```
loss_train = history.history['accuracy']
loss_val = history.history['val_accuracy']
epochs = range(1,21)
plt.plot(epochs, loss_train, 'g', label='Training accuracy')
plt.plot(epochs, loss_val, 'b', label='validation accuracy')
plt.title('Training and Validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



In [42]:

```
score = model.evaluate(X_test, y_test, verbose=1)
print("%s: %.2f%%" % (model.metrics_names[1], score[1]*100))
```

```
132/132 [=====] - 1s 5ms/step - loss: 0.723
1 - accuracy: 0.7302
accuracy: 73.02%
```

In [43]:

```
y_pred = model.predict(X_test)
y_pred = (y_pred > 0.5)
```

In [44]:

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.78	0.76	0.77	1529
1	0.66	0.64	0.65	923
2	1.00	0.00	0.01	275
3	0.82	0.77	0.80	1473
micro avg	0.77	0.69	0.72	4200
macro avg	0.82	0.54	0.56	4200
weighted avg	0.78	0.69	0.70	4200
samples avg	0.69	0.69	0.69	4200

function for creating confusion matrix

In [45]:

```
def print_confusion_matrix(confusion_matrix, class_names, title='Confusion matrix',
                           figsize=(6,6), fontsize=14):
    df_cm = pd.DataFrame(
        confusion_matrix, index=class_names, columns=class_names,
    )
    fig = plt.figure(figsize=figsize)
    try:
        heatmap = sns.heatmap(df_cm, annot=True, fmt="d")
    except ValueError:
        raise ValueError("Confusion matrix values must be integers.")
    heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right',
                                 fontsize=fontsize)
    heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=45, ha='right',
                                 fontsize=fontsize)
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.title(title, fontsize=20)

    return fig
```

printing the confusion matrix

In [46]:

```
from sklearn.metrics import multilabel_confusion_matrix

cnf_matrix = multilabel_confusion_matrix(y_test, y_pred).reshape(4*1, -1)
classes = [str(x) for x in list(y_test.columns.values.tolist())]

print_confusion_matrix(cnf_matrix, classes);
```

